



```
In [ ]: import numpy as np
import pandas as pd
import scipy.stats as stats
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as msno
import plotly.express as px
from sklearn.preprocessing import StandardScaler, MinMaxScaler, PowerTransformer
from sklearn.preprocessing import LabelEncoder
import statsmodels.api as sm
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, precision_score, recall_score
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, auc
from sklearn.metrics import roc_curve
from sklearn.model_selection import train_test_split
from statsmodels.stats.outliers_influence import variance_inflation_factor
from yellowbrick.classifier import ROCAUC
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, ExtraTreesClassifier
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClassifier

plt.rcParams['figure.figsize']=[15,6]

import warnings
warnings.filterwarnings("ignore")
```

```
In [ ]: df=pd.read_csv('train.csv')
pd.set_option('display.max_columns',None)
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

	customer_id	Name	age	gender	security_no	re
0	fffe4300490044003600300030003800	Pattie Morrisey	18	F	XW0DQ7H	
1	fffe43004900440032003100300035003700	Traci Peery	32	F	5K0N3X1	
2	fffe4300490044003100390032003600	Merideth Mcmeen	44	F	1F2TCL3	
3	fffe43004900440036003000330031003600	Eufemia Cardwell	37	M	VJGJ33N	
4	fffe43004900440031003900350030003600	Meghan Kosak	31	F	SVZXCWB	

```
In [ ]: df.shape
        print("Number of rows are",df.shape[0])
        print("Number of columns are",df.shape[1])
```

Number of rows are 36992
Number of columns are 25

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 36992 entries, 0 to 36991
Data columns (total 25 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   customer_id                          36992 non-null  object
1   Name                                 36992 non-null  object
2   age                                  36992 non-null  int64
3   gender                              36992 non-null  object
4   security_no                          36992 non-null  object
5   region_category                      31564 non-null  object
6   membership_category                 36992 non-null  object
7   joining_date                         36992 non-null  object
8   joined_through_referral             36992 non-null  object
9   referral_id                         36992 non-null  object
10  preferred_offer_types                36704 non-null  object
11  medium_of_operation                  36992 non-null  object
12  internet_option                      36992 non-null  object
13  last_visit_time                      36992 non-null  object
14  days_since_last_login                36992 non-null  int64
15  avg_time_spent                       36992 non-null  float64
16  avg_transaction_value                 36992 non-null  float64
17  avg_frequency_login_days             36992 non-null  object
18  points_in_wallet                     33549 non-null  float64
19  used_special_discount                 36992 non-null  object
20  offer_application_preference          36992 non-null  object
21  past_complaint                       36992 non-null  object
22  complaint_status                     36992 non-null  object
23  feedback                             36992 non-null  object
24  churn_risk_score                     36992 non-null  int64
dtypes: float64(3), int64(3), object(19)
memory usage: 7.1+ MB
```

```
In [ ]: df.describe().T
```

```
Out[ ]:
```

	count	mean	std	min	
age	36992.0	37.118161	15.867412	10.000000	23.
days_since_last_login	36992.0	-41.915576	228.819900	-999.000000	8.
avg_time_spent	36992.0	243.472334	398.289149	-2814.109110	60.
avg_transaction_value	36992.0	29271.194003	19444.806226	800.460000	14177.
points_in_wallet	33549.0	686.882199	194.063624	-760.661236	616.
churn_risk_score	36992.0	3.463397	1.409661	-1.000000	3.

```
In [ ]: df.dtypes
```

```
Out[ ]:
```

customer_id	object
Name	object
age	int64
gender	object
security_no	object
region_category	object
membership_category	object
joining_date	object
joined_through_referral	object
referral_id	object
preferred_offer_types	object
medium_of_operation	object
internet_option	object
last_visit_time	object
days_since_last_login	int64
avg_time_spent	float64
avg_transaction_value	float64
avg_frequency_login_days	object
points_in_wallet	float64
used_special_discount	object
offer_application_preference	object
past_complaint	object
complaint_status	object
feedback	object
churn_risk_score	int64
dtype:	object

- joining_date and last_visit_time is in object data type and convert it in to Date data type

Data type Conversion

```
In [ ]: df[['last_visit_time','joining_date']]=df[['last_visit_time','joining_date']].
```

```
In [ ]: df.dtypes
```

```
Out[ ]: customer_id      object
        Name            object
        age             int64
        gender          object
        security_no      object
        region_category  object
        membership_category object
        joining_date      datetime64[ns]
        joined_through_referral object
        referral_id       object
        preferred_offer_types object
        medium_of_operation object
        internet_option   object
        last_visit_time    datetime64[ns]
        days_since_last_login int64
        avg_time_spent     float64
        avg_transaction_value float64
        avg_frequency_login_days object
        points_in_wallet  float64
        used_special_discount object
        offer_application_preference object
        past_complaint     object
        complaint_status  object
        feedback          object
        churn_risk_score   int64
        dtype: object
```


```
In [ ]: df['Joined_Year']=df.joining_date.dt.year
        #df['Joined_Month']=df.joining_date.dt.month_name()
        #df['Joined_day']=df.joining_date.dt.day
        #df['last_visit_Hour']=df.last_visit_time.dt.hour
```

- Extract a new feature 'joined Year' from joined date

Correlation and Covariance

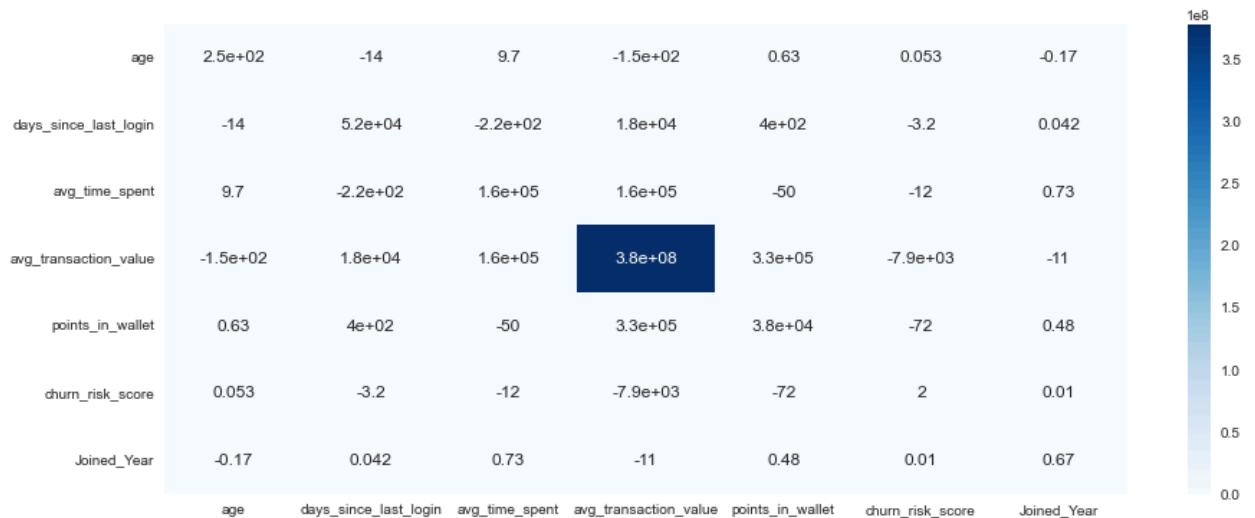
```
In [ ]: sns.heatmap(df.corr(),annot=True,cmap='Blues')
```

```
Out[ ]: <AxesSubplot:>
```



```
In [ ]: sns.heatmap(df.cov(),annot=True,cmap='Blues')
```

```
Out[ ]: <AxesSubplot:>
```



- From the abovemheatmap we can get to known that there is no correlation.

```
In [ ]: df.describe().T
```

```
Out[ ]:
```

	count	mean	std	min	
age	36992.0	37.118161	15.867412	10.000000	23.
days_since_last_login	36992.0	-41.915576	228.819900	-999.000000	8.
avg_time_spent	36992.0	243.472334	398.289149	-2814.109110	60.
avg_transaction_value	36992.0	29271.194003	19444.806226	800.460000	14177.
points_in_wallet	33549.0	686.882199	194.063624	-760.661236	616.
churn_risk_score	36992.0	3.463397	1.409661	-1.000000	3.
Joined_Year	36992.0	2016.006569	0.819384	2015.000000	2015.

It looks like -999 is an 'Error', i.e., the website didn't populate the variable when the data was recorded. Hence replacing it with median as their value so that we can visualize how the data is spread out

```
In [ ]: df.loc[df['days_since_last_login'] < 0, 'days_since_last_login'].value_counts(r
```

```
Out[ ]: -999    1.0
        Name: days_since_last_login, dtype: float64
```

```
In [ ]: df.loc[df['days_since_last_login'] < 0, 'days_since_last_login'] = df['days_sin
```

```
In [ ]: df.loc[df['days_since_last_login'] < 0, 'days_since_last_login']
```

```
Out[ ]: Series([], Name: days_since_last_login, dtype: int64)
```

```
In [ ]: '''for i in df.columns:
        if df[i].dtypes in [np.int64, np.number]:
            df.loc[df[i] < 0, i] = df[i].loc[df[i] > 0].median()'''
```

```
Out[ ]: 'for i in df.columns:\n    if df[i].dtypes in [np.int64, np.number]:\n        df.loc[df[i] < 0, i] = df[i].loc[df[i] > 0].median()'
```

Missing_values

```
In [ ]: # Assuming df is your DataFrame
        missing_values = df.isnull().sum()
        print(missing_values)
```

```

-----
AttributeError                                Traceback (most recent call last)
Input In [167], in <cell line: 1>()
----> 1 df_ = df[df.loc[:, 'str'].contains('?')]
      2 df_

File ~\anaconda3\lib\site-packages\pandas\core\generic.py:5575, in NDFrame.__getattr__(self, name)
    5568 if (
    5569     name not in self._internal_names_set
    5570     and name not in self._metadata
    5571     and name not in self._accessors
    5572     and self._info_axis._can_hold_identifiers_and_holds_name(name)
    5573 ):
    5574     return self[name]
-> 5575 return object.__getattr__(self, name)

AttributeError: 'DataFrame' object has no attribute 'str'


```

```
In [ ]: df.replace({'?':np.nan },inplace=True)
```

- Replace '?' with null values

```
In [ ]: msno.bar(df)
```

```
Out[ ]: <AxesSubplot:>
```



```
In [ ]: df.isnull().sum()[df.isnull().sum()!=0]
```

```
Out[ ]: region_category      5428
        joined_through_referral  5438
        preferred_offer_types    288
        medium_of_operation      5393
        points_in_wallet        3443
        dtype: int64
```

From this we observed that the null values present in the columns

- region_category
- point_in_wallet
- joined_through_referral
- preferred_offer_types
- medium_of_operation

```
In [ ]: Null_values=(df.isnull().sum()/len(df)*100)[(df.isnull().sum()/len(df)*100)!=0]
        print(Null_values)
```

```
region_category      14.673443
joined_through_referral  14.700476
preferred_offer_types    0.778547
medium_of_operation      14.578828
points_in_wallet        9.307418
dtype: float64
```

Null_Values Handling

- We consider to drop the rows having null values less than 5%
- For Categorical Columns we are imputing with Mode
- For Numerical columns we are imputing with Median


```
In [ ]: for i in Null_values.index:
        if Null_values[i][Null_values[i]<=5]:
            df.dropna(subset=[i],axis=0,inplace=True)
        elif Null_values[i][Null_values[i]>5]:
            if df[i].dtypes in [np.int64, np.number]:
                df[i].fillna(df[i].median(),inplace=True)
            elif df[i].dtypes == np.object_:
                df[i].fillna(df[i].mode()[0],inplace=True)
```

```
In [ ]: Null_values=(df.isnull().sum()/len(df)*100)[(df.isnull().sum()/len(df)*100)!=0]
        print(Null_values)
```

```
Series([], dtype: float64)
```

```
In [ ]: sns.countplot(df.churn_risk_score)
```

```
Out[ ]: <AxesSubplot:xlabel='churn_risk_score', ylabel='count'>
```

Feature Engineering

```
In [ ]: pd.crosstab(df.churn_risk_score,df.feedback)
```

Out[]:

	feedback	No reason specified	Poor Customer Service	Poor Product Quality	Poor Website	Products always in Stock	Quality Customer Care
churn_risk_score							
-1		214	195	197	208	36	40
1		0	0	0	0	675	626
2		0	0	0	0	660	688
3		2062	2041	2015	2080	0	0
4		1977	2024	2100	2047	0	0
5		1981	1935	1992	1891	0	0

Churn risk rate -1 is not feasible value so we have done feature engineering to impute the value Compare the feedback and assign accordingly

```
In [ ]: def feed(x, y):
    l1 = ['Poor Quality','Too many ads','Poor Website','Poor Customer Service']
    if y == -1:
        if x in l1:
            return 5
        else:
            return 1
    else:
        return y

df["churn_risk_score"] = df.apply(lambda x: feed(x['feedback'],x['churn_risk_s
```

```
In [ ]: df.churn_risk_score.value_counts()
```

```
Out[ ]: 5    10341
        3    10339
        4    10098
        1     3207
        2     2719
        Name: churn_risk_score, dtype: int64
```

Labelling

- We have Churn risk rate from 1 to 5
- we are bucketing the label in order for better prediction
- risk rate 1 and 2 are low risk so we assign those to label 0
- risk rate 3 falls in both low and high, so we keep it as a standalone label 1
- risk rate 4 and 5 are high risk rate so we assing it to label 2


```
In [ ]: def bucket(x):
        if (x == 1) | (x == 2):
            return 0
        elif (x == 3):
            return 1
        else:
            return 2
        df["churn_risk_score"] = df.apply(lambda x: bucket(x['churn_risk_score']), axis=1)
```

```
In [ ]: df.churn_risk_score.value_counts()
```

```
Out[ ]: 2    20439
        1    10339
        0     5926
        Name: churn_risk_score, dtype: int64
```

```
In [ ]: sns.countplot(df.churn_risk_score)
```

```
Out[ ]: <AxesSubplot:xlabel='churn_risk_score', ylabel='count'>
```



Dropping In-significant variables

```
In [ ]: df.avg_frequency_login_days.value_counts()
```

```
Out[ ]: Error          3496
13.0          1382
19.0          1351
8.0           1350
14.0          1349
...
28.191570401129514      1
41.73357294995208      1
-11.515939810499656     1
45.71683637272365      1
27.8399274405269       1
Name: avg_frequency_login_days, Length: 1632, dtype: int64
```

Looks like avg_freq_login_days(Represents the no. of times a customer has logged in to the website) variable is holding numeric datatype. Hence converted to float.

ERROR value infers that the website was unable to register the avg_freq_login_days. It could be due to various factors like software glitches, etc. Also, the variable days since last login and average frequency login days holds redundancy in terms of their usage. Hence dropping the variable.

Customer-id, Name, security_no are unique variables. referral_id is completely irrelevant to the dataset. Hence, dropping the above mentioned variables.

```
In [ ]: df=df.drop(columns=['customer_id','Name','security_no','referral_id','avg_freq
```

```
In [ ]: df['churn_risk_score'] = df['churn_risk_score'].astype('object')
```

Univariate_analysis

```
In [ ]: Caterogical_columns=df.select_dtypes(include=np.object_).columns
print(Caterogical_columns)

Index(['gender', 'region_category', 'membership_category',
      'joined_through_referral', 'preferred_offer_types',
      'medium_of_operation', 'internet_option', 'used_special_discount',
      'offer_application_preference', 'past_complaint', 'complaint_status',
      'feedback', 'churn_risk_score'],
      dtype='object')
```

```
In [ ]: Numerical_columns=df.select_dtypes(include=np.number).columns
print(Numerical_columns)
len(Numerical_columns)

Index(['age', 'days_since_last_login', 'avg_time_spent',
      'avg_transaction_value', 'points_in_wallet', 'Joined_Year'],
      dtype='object')
```

Out[]: 6



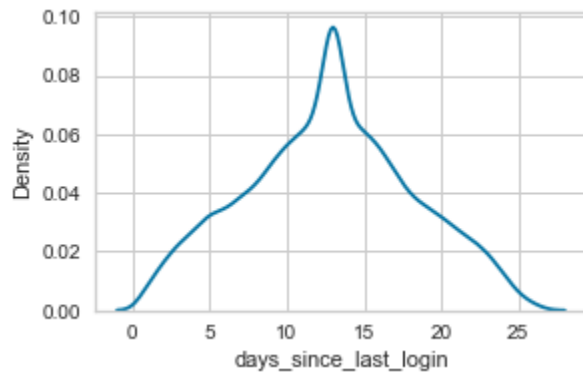

```
In [ ]: df.dtypes
```


```
Out[ ]: age                int64
gender                object
region_category       object
membership_category   object
joined_through_referral object
preferred_offer_types  object
medium_of_operation    object
internet_option        object
days_since_last_login int64
avg_time_spent         float64
avg_transaction_value  float64
points_in_wallet       float64
used_special_discount  object
offer_application_preference object
past_complaint         object
complaint_status      object
feedback              object
churn_risk_score       object
Joined_Year            int64
dtype: object
```

```
In [ ]: nrows=2
ncols=3
iterator=1

for i in Numerical_columns:
    plt.subplot(nrows,ncols,iterator)
    sns.kdeplot(df.loc[:,i])
    plt.show()
    iterator+=1

plt.tight_layout()
plt.show()
```











<Figure size 1080x432 with 0 Axes>

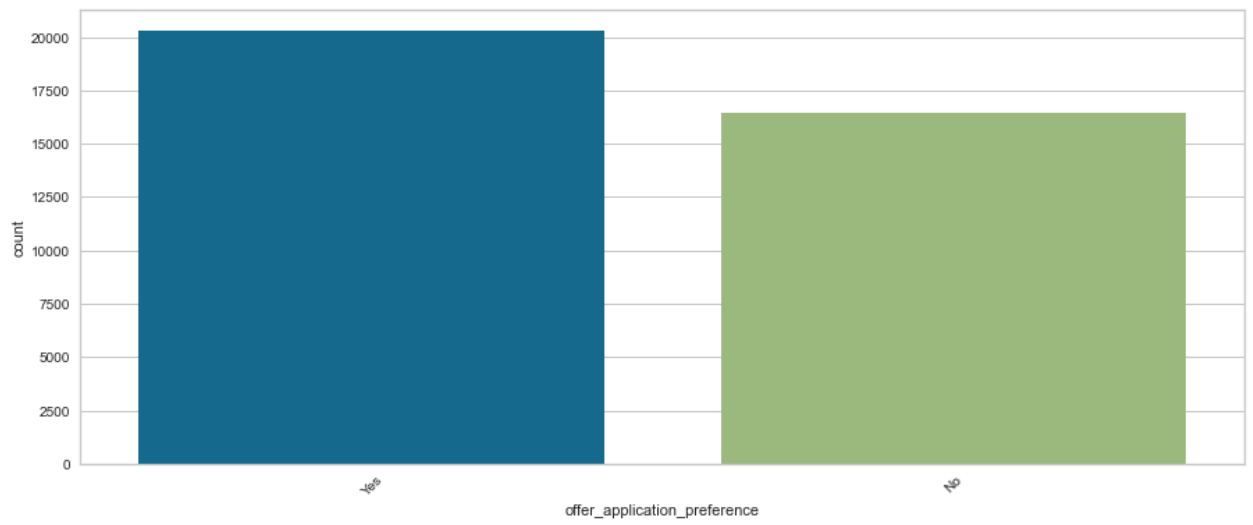


From the above figure we can get to know about the distribution of the data for all the numerical columns.




```
In [ ]: nrow=9
ncol=2
iterator=1


for i in Caterogical_columns:
    #plt.subplot(nrow,ncol,iterator)
    sns.countplot(df.loc[:,i])
    #iterator+=1
    plt.xticks(rotation=45)
    plt.show()
```











The above plots represents the data available in categorical columns and their distribution.

Inference

- The Age category shows that the Age is widely spread from 10-64 with almost equal weightage.
- The Gender is almost equally spread in the data except the unknown subclass
- The region category Town is having the maximum counts. And the region Village is having the minimum counts. The plot shows that the town population is attracted to this particular e-commerce site.
- The Membership that the Basic Membership and No Membership are having the highest count. With the platinum membership being the lowest in count
- The preferred_offer_types is almost equally spread in the data except the missing values. We cannot impute missing values as this the variable is related to personal information.
- From the Barplot, we can infer that both Desktop and Smartphone is spread equally. We can also infer that only 10% of people using both Smartphone and Desktop.
- The Internet option is showing equal weightage to all the subclasses being Wi-Fi, Mobile, Fiber_Optic.
- The days_since_last_login variable is holding the number of days since the customer has logged in. The plot shows that the average lies around 13. The maximum days since logged-in is 31. And the minimum is 1.
- Maximum number of negative feedback for the variable is poor product


quality and maximum number of positive feedback for the variable is user friendly website and reasonable price.




- The complaint status' subclasses holds almost equal weights except 'Not Applicable'.
- In Medium of operation Desktop users are more than smartphone
- Used special discount column has shos that most people are not using special discount offers
- Point in wallet has the maximum value of 500 to 800



Bivariate analysis


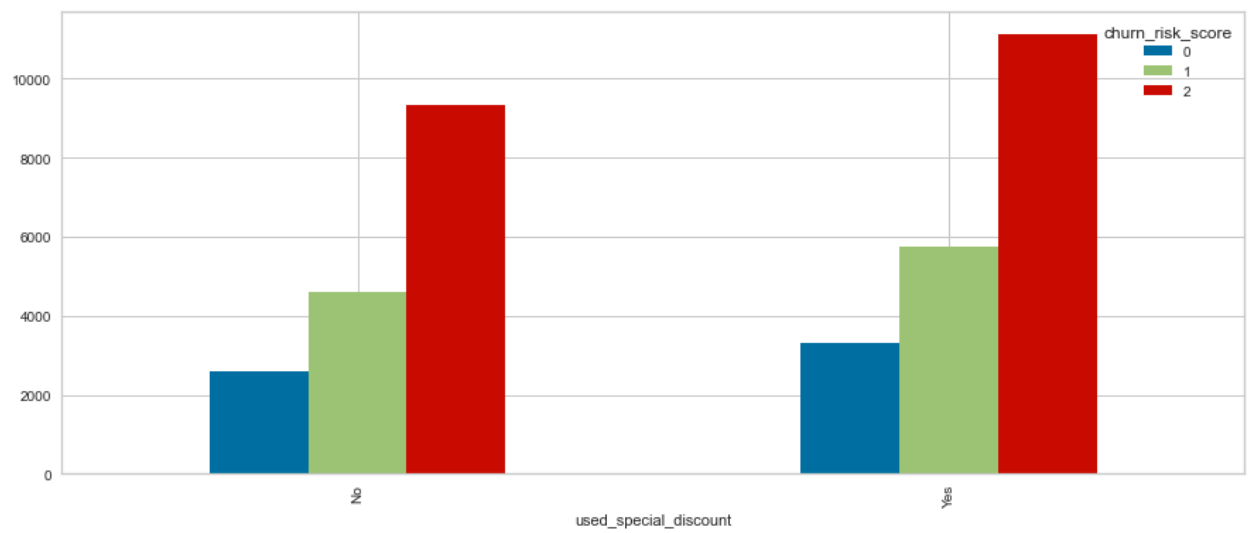

```
In [ ]: nrows=9
ncols=2
iterator=1



for i in Caterogical_columns:
    #plt.subplot(nrows,ncols,iterator)
    pd.crosstab(df.loc[:,i],df.churn_risk_score).plot(kind='bar')
    #sns.boxplot(x=df.loc[:,i],y=df.churn_risk_score)
    #iterator+=1
    plt.show()
```














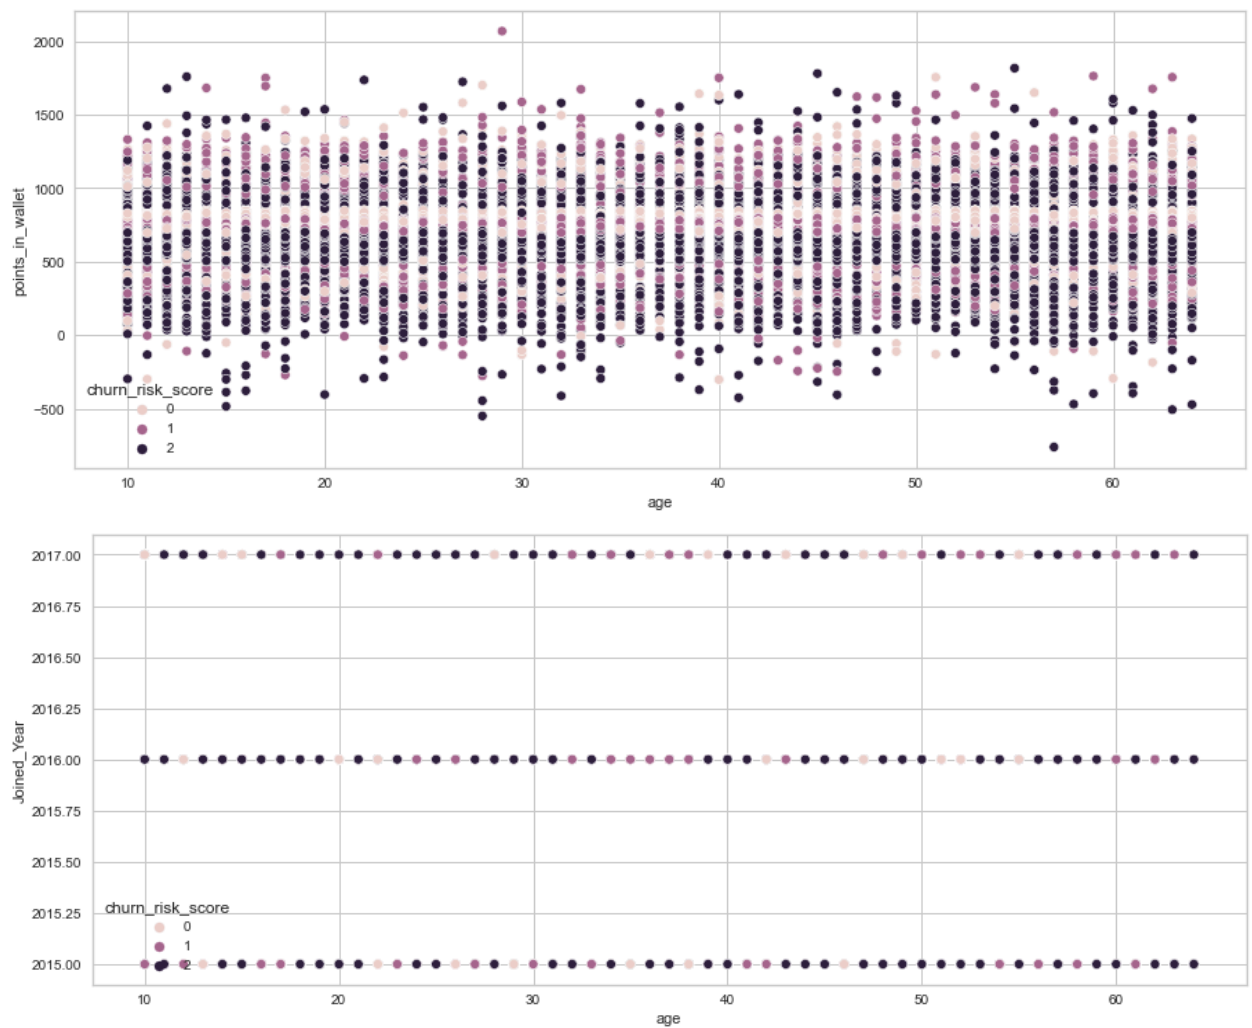




The above plot represents the relationship between the target variable and the other categorical variables.

```
In [ ]: nrow=9
ncol=2
iterator=1
for i in Numerical_columns:
    if i!='age':
        sns.scatterplot(x=df.age,y=df.loc[:,i],hue=df.churn_risk_score)
        plt.show()
```





The Above plot represents the relationship between the target variable and the other numerical variables.

Inference

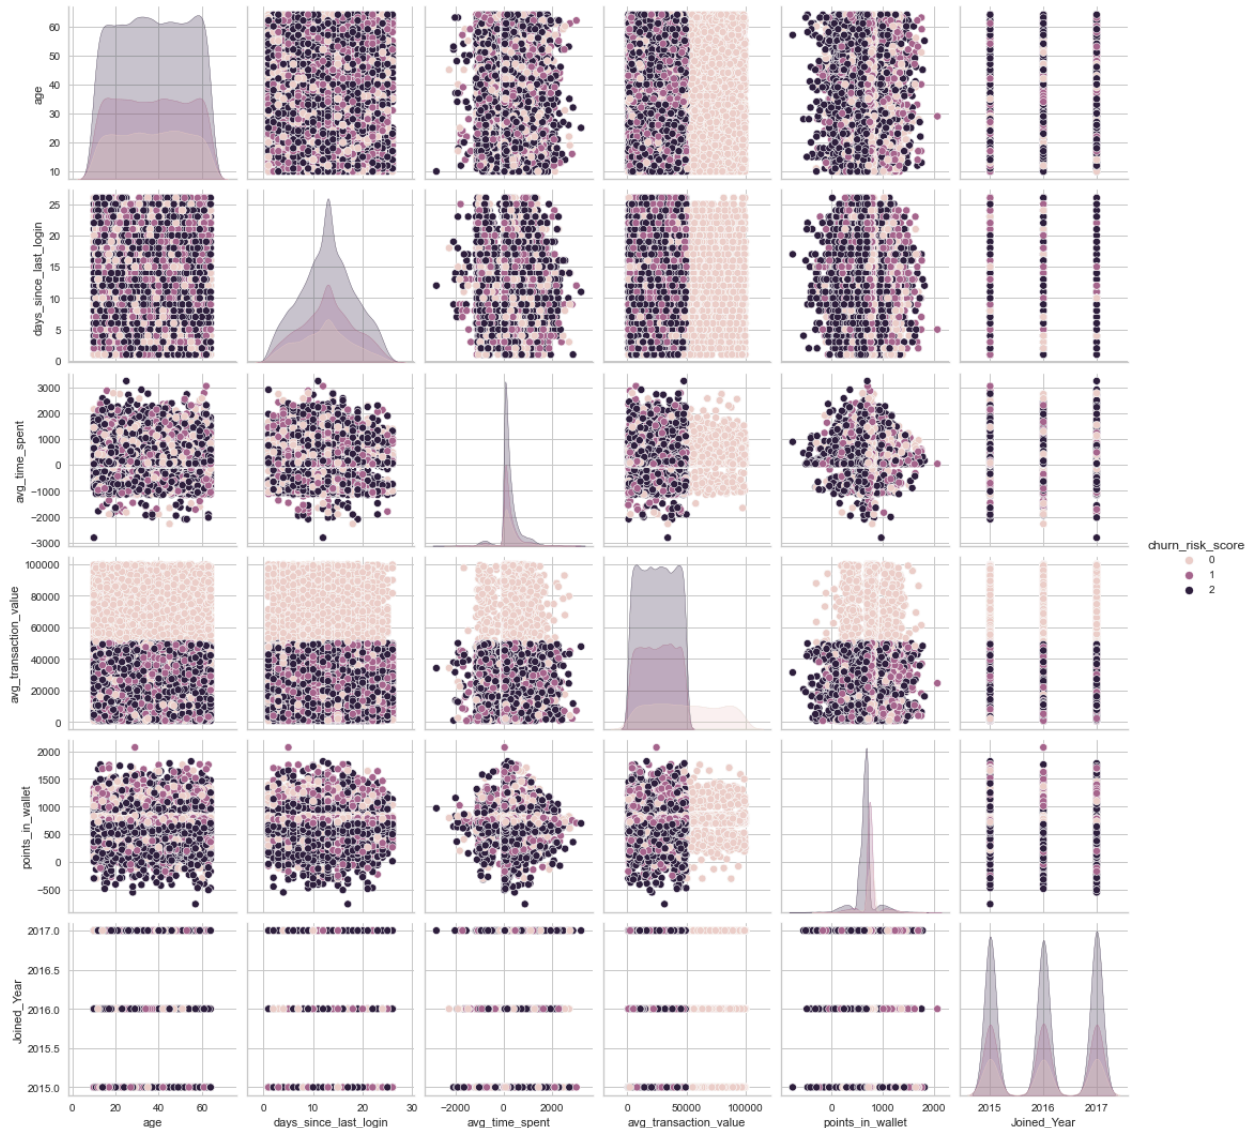
- In Gender Both values holds equal weightage for churn risk rate
- In Region category Churn risk rate is high for town
- Basic and No membership having high churn risk rate
- Platinum and premium holds low churn risk rate
- Silver and Gold holds both low as well as high
- Without offers customers having high churn score whereas customer using coupons having low churn risk
- Complaint Status Not applicable holds high churn score
- Positive feedback having low churn score whereas Negative feedback having high churn score
- Customers with average transaction value greater than 50000 holds

low churn score

```
In [ ]: plt.figure(figsize=(20,20))
sns.pairplot(df, hue='churn_risk_score')
```

```
Out[ ]: <seaborn.axisgrid.PairGrid at 0x1bd4f018fd0>
```

<Figure size 1440x1440 with 0 Axes>



The average transaction value is holding maximum range for churn risk score 0 irrespective of the age, days_since_last_login, avg_time_Spent, points_in_wallet. And the values are low for churn risk score 1,2. We are able to see the separation of clusters between 0 and 1,2 after bucketising. The points_in_wallet shows the dominance of cluster 0 for points above 500, whereas cluster 2 is showing its dominance for points below 500.

```
In [ ]: for col in Categorical_columns:
print("-----"+str(col)+"-----")
```

```
print("unique Values :",df[col].unique()) # to print categories name only
print("Value counts of unique values :\n",df[col].value_counts()) # to pri
print("-----")
```

```

-----gender-----
unique Values : ['F' 'M' 'Unknown']
Value counts of unique values :
  F          18348
  M          18298
  Unknown      58
Name: gender, dtype: int64
-----

-----region_category-----
unique Values : ['Village' 'City' 'Town']
Value counts of unique values :
  Town          19404
  City          12635
  Village        4665
Name: region_category, dtype: int64
-----

-----membership_category-----
unique Values : ['Platinum Membership' 'Premium Membership' 'No Membership'
'Gold Membership' 'Silver Membership' 'Basic Membership']
Value counts of unique values :
  Basic Membership      7662
  No Membership        7632
  Gold Membership      6742
  Silver Membership    5935
  Premium Membership   4427
  Platinum Membership  4306
Name: membership_category, dtype: int64
-----

-----joined_through_referral-----
unique Values : ['No' 'Yes']
Value counts of unique values :
  No      21126
  Yes     15578
Name: joined_through_referral, dtype: int64
-----

-----preferred_offer_types-----
unique Values : ['Gift Vouchers/Coupons' 'Credit/Debit Card Offers' 'Without Of
fers']
Value counts of unique values :
  Gift Vouchers/Coupons      12349
  Credit/Debit Card Offers   12274
  Without Offers             12081
Name: preferred_offer_types, dtype: int64
-----

-----medium_of_operation-----
unique Values : ['Desktop' 'Smartphone' 'Both']
Value counts of unique values :
  Desktop      19154
  Smartphone   13766

```

```

Both          3784
Name: medium_of_operation, dtype: int64
-----
-----internet_option-----
unique Values : ['Wi-Fi' 'Mobile_Data' 'Fiber_Optic']
Value counts of unique values :
  Wi-Fi          12310
Mobile_Data     12247
Fiber_Optic     12147
Name: internet_option, dtype: int64
-----
-----used_special_discount-----
unique Values : ['Yes' 'No']
Value counts of unique values :
  Yes      20182
No       16522
Name: used_special_discount, dtype: int64
-----
-----offer_application_preference-----
unique Values : ['Yes' 'No']
Value counts of unique values :
  Yes      20282
No       16422
Name: offer_application_preference, dtype: int64
-----
-----past_complaint-----
unique Values : ['No' 'Yes']
Value counts of unique values :
  No      18446
Yes      18258
Name: past_complaint, dtype: int64
-----
-----complaint_status-----
unique Values : ['Not Applicable' 'Solved' 'Solved in Follow-up' 'Unsolved'
'No Information Available']
Value counts of unique values :
  Not Applicable      18446
Unsolved              4615
Solved                4579
Solved in Follow-up   4542
No Information Available 4522
Name: complaint_status, dtype: int64
-----
-----feedback-----
unique Values : ['Products always in Stock' 'Quality Customer Care' 'Poor Website'
'No reason specified' 'Poor Product Quality' 'Poor Customer Service']

```

```

'Too many ads' 'User Friendly Website' 'Reasonable Price']
Value counts of unique values :
Poor Product Quality      6304
No reason specified       6234
Too many ads              6230
Poor Website              6226
Poor Customer Service     6195
Reasonable Price          1408
User Friendly Website     1382
Products always in Stock  1371
Quality Customer Care     1354
Name: feedback, dtype: int64
-----
-----
-----churn_risk_score-----
unique Values : [0 2 1]
Value counts of unique values :
 2    20439
1    10339
0     5926
Name: churn_risk_score, dtype: int64
-----
-----

```

```
In [ ]: #df1.age.describe()# Have to split age in to Teen = 13-19 yrs. Adult = 20-39 y
```

```
def age_categorize(age):
```

```

    return('Child')
elif (age<=19):
    return('Teen')
elif (age<=39):
    return('Adult')
elif (age<=59):
    return('Middle_Age_Adult')
else:
    return('Senior_Adult')
```

```
df['age_category']=df.age.apply(age_categorize)
```

Skewness

```
In [ ]: df.skew()
```



```
Out[ ]: age                -0.007368
        days_since_last_login  0.021134
        avg_time_spent        0.538800
        avg_transaction_value  1.009753
        points_in_wallet      -0.102518
        churn_risk_score      -0.790561
        Joined_Year          -0.011602
        dtype: float64
```

Data in various columns are positively as well as negatively skewed

Outlier

```
In [ ]: q1=df.quantile(.25)
        q3=df.quantile(.75)
        IQR=q3-q1
        ll=q1-1.5*IQR
        ul=q3+1.5*IQR
        wt_outliers=df.loc[((df>ul)|(df<ll)).any(axis=1)]
        wt_outliers.shape
```

```
Out[ ]: (9620, 19)
```

We are having outliers in our data of around 9620 rows but we keep our outliers in our data

Statistical Test

Null hypothesis (Ho) : Predictor and Target are Independent Alternate hypothesis (Ha) : Predictor and Target are Dependent Confidence Interval : 0.95 level_of_significance : 0.05

```
In [ ]: from scipy.stats import shapiro,levne,contingency,chisquare,ttest_ind,f_oneway
```

```
In [ ]: num_cols=df.select_dtypes(include=np.number).columns
        cat_columns=df.select_dtypes(include=np.object_).columns
```

```
In [ ]: num_cols
```

```
Out[ ]: Index(['age', 'days_since_last_login', 'avg_time_spent',
              'avg_transaction_value', 'points_in_wallet', 'Joined_Year'],
              dtype='object')
```

```
In [ ]: import scipy.stats as stats
```

```
signif_feats1=[]
test_stats1=[]
p_value1=[]
signif_feats2=[]
```

```

test_stats2=[]
p_value2=[]
for i in num_cols:
    one=df.loc[df.churn_risk_score==0,i]
    two=df.loc[df.churn_risk_score==1,i]
    three=df.loc[df.churn_risk_score==2,i]
    teststats,pvalue=stats.f_oneway(one,two,three)
    if pvalue <0.05:
        signif_feats1.append(i)
        test_stats1.append(teststats)
        p_value1.append(pvalue)
    else:
        signif_feats2.append(i)
        test_stats2.append(teststats)
        p_value2.append(pvalue)

for i in cat_columns:
    if i!='churn_risk_score':
        test_stat, pvalue, dof, expected_value = chi2_contingency(pd.crosstab(
        if pvalue <0.05:
            signif_feats1.append(i)
            test_stats1.append(teststats)
            p_value1.append(pvalue)
        else:
            signif_feats2.append(i)
            test_stats2.append(teststats)
            p_value2.append(pvalue)

Dependent_Features=pd.DataFrame({'Features':signif_feats1,'Test_Statistics':te
Independent_Features=pd.DataFrame({'Features':signif_feats2,'Test_Statistics':

```

```
In [ ]: print(Dependent_Features)
```

	Features	Test_Statistics	PValue
0	days_since_last_login	34.385545	1.203709e-15
1	avg_time_spent	24.024012	3.743933e-11
2	avg_transaction_value	4749.020578	0.000000e+00
3	points_in_wallet	1781.980356	0.000000e+00
4	region_category	0.783061	2.894598e-18
5	membership_category	0.783061	0.000000e+00
6	joined_through_referral	0.783061	1.832495e-26
7	preferred_offer_types	0.783061	5.892588e-64
8	medium_of_operation	0.783061	6.357438e-17
9	used_special_discount	0.783061	4.818971e-02
10	offer_application_preference	0.783061	2.566210e-20
11	past_complaint	0.783061	4.036414e-02
12	feedback	0.783061	0.000000e+00

```
In [ ]: print(Independent_Features)
```

	Features	Test_Statistics	PValue
0	age	0.702638	0.495284
1	Joined_Year	0.783061	0.457013
2	gender	0.783061	0.606347
3	internet_option	0.783061	0.314392
4	complaint_status	0.783061	0.195914

```
In [ ]: df_new=df.drop(columns=['age','gender','internet_option','complaint_status'],'J
```

```
In [ ]: df_new.shape
```

```
Out[ ]: (36704, 14)
```

Scaling

```
In [ ]: df1=df_new.select_dtypes(include=np.number)
ss=StandardScaler()
df_s=ss.fit_transform(df1)
df_s=pd.DataFrame(df_s,columns=df1.columns,index=df1.index)
df_s.head()

df_s
```

```
Out[ ]:
```

	days_since_last_login	avg_time_spent	avg_transaction_value	points_in_v
0	0.778932	0.143803	1.220169	0.50
1	0.594431	0.158149	-0.845189	0.00
2	0.225429	0.685315	-0.424135	-1.00
3	-0.328075	-0.477681	-0.207527	-0.60
4	1.332435	-0.327285	-0.246395	-0.10
...
36987	-1.988585	-2.246341	-0.102728	-0.20
36988	0.040928	-2.214786	-0.936133	-0.80
36989	-0.143573	-0.222239	0.455167	-0.00
36990	0.409930	0.601022	-1.383012	-2.60
36991	0.409930	-0.412583	-1.392740	0.10

36704 rows × 4 columns

```
In [ ]: df_kk=pd.concat([df_s,df.churn_risk_score],axis=1)
df_kk['churn_risk_score']=df_kk.churn_risk_score.astype('int')
```

Encoding

```
In [ ]: df_cat = df.select_dtypes(include=[np.object])
df_cat=df_cat.drop(['churn_risk_score'],axis=1)
for i in df_cat.columns:
    df_cat[i]=LabelEncoder().fit_transform(df_cat[i])
df_cat.head()
```

```
Out[ ]:
```

	gender	region_category	membership_category	joined_through_referral	pref
0	0	2	3	0	
1	0	0	4	0	
2	0	1	2	1	
3	1	0	2	1	
4	0	0	2	0	

```
In [ ]: #df_cat=pd.get_dummies(df_new,drop_first=True)
#df_cat.drop(columns=['joining_date','last_visit_time'],inplace=True)
```

```
In [ ]: df_new1=pd.concat([df_kk,df_cat],axis=1)
df_new1
```

```
Out[ ]:
```

	days_since_last_login	avg_time_spent	avg_transaction_value	points_in_v
0	0.778932	0.143803	1.220169	0.50
1	0.594431	0.158149	-0.845189	0.00
2	0.225429	0.685315	-0.424135	-1.00
3	-0.328075	-0.477681	-0.207527	-0.60
4	1.332435	-0.327285	-0.246395	-0.10
...
36987	-1.988585	-2.246341	-0.102728	-0.20
36988	0.040928	-2.214786	-0.936133	-0.80
36989	-0.143573	-0.222239	0.455167	-0.00
36990	0.409930	0.601022	-1.383012	-2.60
36991	0.409930	-0.412583	-1.392740	0.10

36704 rows × 17 columns

```
In [ ]: X=df_new1.drop('churn_risk_score',1)
y=df.churn_risk_score.astype('int')
```

```
In [ ]: y.dtype
```

```
Out[ ]: dtype('int32')
```

```
In [ ]: print(X.shape,y.shape)
```

```
(36704, 16) (36704,)
```

```
In [ ]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(29363, 16)
```

```
(7341, 16)
```

```
(29363,)
```

```
(7341,)
```

```
In [ ]: X.dtypes
```

```
Out[ ]: days_since_last_login      float64
avg_time_spent                    float64
avg_transaction_value             float64
points_in_wallet                  float64
gender                           int32
region_category                   int32
membership_category               int32
joined_through_referral           int32
preferred_offer_types             int32
medium_of_operation               int32
internet_option                   int32
used_special_discount             int32
offer_application_preference       int32
past_complaint                    int32
complaint_status                  int32
feedback                          int32
dtype: object
```

```
In [ ]: vif=[]
for i in range (0,df_s.shape[1]):
    vif.append(variance_inflation_factor(df_s.values,i))
pd.DataFrame({'features':df_s.columns,'VIF':vif})
```

```
Out[ ]:
```

	features	VIF
0	days_since_last_login	1.008061
1	avg_time_spent	1.008226
2	avg_transaction_value	1.007821
3	points_in_wallet	1.007175

VIF score for all values are less than 10,hence there is no multicollinearity

Model Building

List of Models:

- LogisticRegression
- DecisionTreeClassifier
- RandomForestClassifier
- ExtraTreesClassifier
- XGBClassifier
- LGBMClassifier
- AdaBoostClassifier
- GradientBoostingClassifier

```
In [ ]: score_card = pd.DataFrame(columns=['Model', 'Precision Score', 'Recall Score',  
                                         'False Negatives', 'Kappa Score', 'f1-score'])
```

```
In [ ]: def update_score_card(model, FN_values, model_name):  
    y_pred = model.predict(X_test)  
    global score_card  
    score_card = score_card.append({'Model': model_name,  
                                    'Precision Score': precision_score(y_test,  
                                'Recall Score': recall_score(y_test, y_pred,  
                                'False Negatives': FN_values,  
                                'Kappa Score': cohen_kappa_score(y_test, y_pred,  
                                'f1-score': f1_score(y_test, y_pred, average='weighted',  
                                ignore_index = True)
```

```
In [ ]: from sklearn.linear_model import LogisticRegression  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.ensemble import RandomForestClassifier,ExtraTreesClassifier  
from xgboost import XGBClassifier  
from lightgbm import LGBMClassifier  
from sklearn.ensemble import AdaBoostClassifier,GradientBoostingClassifier
```

```
In [ ]: LR = LogisticRegression(multi_class='multinomial', solver='lbfgs')  
LR_Model=LR.fit(X_train,y_train)  
y_pred_xtest=LR_Model.predict(X_test)  
print(classification_report(y_test,y_pred_xtest))
```

	precision	recall	f1-score	support
0	0.69	0.58	0.63	1185
1	0.59	0.47	0.52	2044
2	0.75	0.85	0.79	4112
accuracy			0.70	7341
macro avg	0.67	0.63	0.65	7341
weighted avg	0.69	0.70	0.69	7341

```
In [ ]: print(accuracy_score(y_test,y_pred_xtest))
```


```
0.7011306361531127
```

```
In [ ]: print(confusion_matrix(y_test,y_pred_xtest))
```

```
[[ 686  201  298]
 [ 181  967  896]
 [ 134  484 3494]]
```

```
In [ ]: #For logistic regression model, the roc curve with yellowbrick package
LR_visualizer = ROCAUC(LR_Model)
```

```
LR_visualizer.fit(X_train, y_train)      # Fit the training data to the visu
LR_visualizer.score(X_test, y_test)      # Evaluate the model on the test da
LR_visualizer.show()
```




```
Out[ ]: <AxesSubplot:title={'center':'ROC Curves for LogisticRegression'}, xlabel='False Positive Rate', ylabel='True Positive Rate'>
```

```
In [ ]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, auc, mul
```

```
In [ ]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred_xtest, cmap='Reds')
```

```
Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1bd52f0b280>
```



```
In [ ]: LR_mul = multilabel_confusion_matrix(y_test, y_pred_xtest)
LR_mul
```

```
Out[ ]: array([[5841, 315],
               [ 499, 686]],

            [[4612, 685],
             [1077, 967]],

            [[2035, 1194],
             [ 618, 3494]]], dtype=int64)
```

```
In [ ]: LR_FN = LR_mul[2][1][0]
update_score_card(LR_Model, LR_FN, 'Logistic Regression')
```

```
In [ ]:
```

```
In [ ]:
```

Decision_Tree

```
In [ ]: dt=DecisionTreeClassifier()
DT_Model=dt.fit(X_train,y_train)
y_pred_xtest=DT_Model.predict(X_test)
print(classification_report(y_test,y_pred_xtest))
```



	precision	recall	f1-score	support
0	0.90	0.92	0.91	1185
1	0.88	0.86	0.87	2044
2	0.92	0.93	0.93	4112
accuracy			0.91	7341
macro avg	0.90	0.90	0.90	7341
weighted avg	0.91	0.91	0.91	7341

```
In [ ]: print(accuracy_score(y_test,y_pred_xtest))
        print(confusion_matrix(y_test,y_pred_xtest))
```

```
0.9091404440811879
[[1093  27  65]
 [ 40 1759 245]
 [ 75  215 3822]]
```

```
In [ ]: LR_visualizer = ROCAUC(DT_Model)


        LR_visualizer.fit(X_train, y_train)           # Fit the training data to the visu
        LR_visualizer.score(X_test, y_test)           # Evaluate the model on the test da
        LR_visualizer.show()
```



```
Out[ ]: <AxesSubplot:title={'center':'ROC Curves for DecisionTreeClassifier'}, xlabel='False Positive Rate', ylabel='True Positive Rate'>
```

```
In [ ]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred_xtest, cmap='Reds')
```

```
Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1bd55eb7d60>
```



```
In [ ]: LR_mul = multilabel_confusion_matrix(y_test, y_pred_xtest)
LR_mul
```

```
Out[ ]: array([[6041, 115],
               [ 92, 1093]],

               [[5055, 242],
                [ 285, 1759]],

               [[2919, 310],
                [ 290, 3822]]], dtype=int64)
```

```
In [ ]: LR_FN = LR_mul[2][1][0]
update_score_card(DT_Model, LR_FN, 'Decision Tree')
```

RandomForest

```
In [ ]: rf=RandomForestClassifier()
RF_Model=rf.fit(X_train,y_train)
y_pred_xtest=RF_Model.predict(X_test)
print(classification_report(y_test,y_pred_xtest))
```


	precision	recall	f1-score	support
0	1.00	0.92	0.96	1185
1	0.88	0.90	0.89	2044
2	0.94	0.95	0.94	4112
accuracy			0.93	7341
macro avg	0.94	0.92	0.93	7341
weighted avg	0.93	0.93	0.93	7341

```
In [ ]: print(accuracy_score(y_test,y_pred_xtest))
        print(confusion_matrix(y_test,y_pred_xtest))
```

```
0.9317531671434409
[[1094  31  60]
 [  0 1845 199]
 [  0  211 3901]]
```

```
In [ ]: LR_visualizer = ROCAUC(RF_Model)


        LR_visualizer.fit(X_train, y_train)           # Fit the training data to the visu
        LR_visualizer.score(X_test, y_test)           # Evaluate the model on the test da
        LR_visualizer.show()
```



```
Out[ ]: <AxesSubplot:title={'center':'ROC Curves for RandomForestClassifier'}, xlabel='False Positive Rate', ylabel='True Positive Rate'>
```

```
In [ ]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred_xtest, cmap='Reds')
```

```
Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1bd59151700>
```



```
In [ ]: LR_mul = multilabel_confusion_matrix(y_test, y_pred_xtest)
LR_mul
```

```
Out[ ]: array([[6156,  0],
               [ 91, 1094]],

               [[5055, 242],
                [199, 1845]],

               [[2970, 259],
                [211, 3901]]], dtype=int64)
```

```
In [ ]: LR_FN = LR_mul[2][1][0]
update_score_card(RF_Model, LR_FN, 'Random Forest')
```

ExtraTreesClassifier

```
In [ ]: et=ExtraTreesClassifier()
ET_Model=et.fit(X_train,y_train)
y_pred_xtest=ET_Model.predict(X_test)
print(classification_report(y_test,y_pred_xtest))
```


	precision	recall	f1-score	support
0	1.00	0.92	0.96	1185
1	0.86	0.86	0.86	2044
2	0.92	0.94	0.93	4112
accuracy			0.91	7341
macro avg	0.93	0.91	0.92	7341
weighted avg	0.92	0.91	0.92	7341

```
In [ ]: print(accuracy_score(y_test,y_pred_xtest))
        print(confusion_matrix(y_test,y_pred_xtest))
```

```
0.914997956681651
[[1096  27  62]
 [   0 1765 279]
 [   1  255 3856]]
```

```
In [ ]: LR_visualizer = ROCAUC(ET_Model)


        LR_visualizer.fit(X_train, y_train)           # Fit the training data to the visu
        LR_visualizer.score(X_test, y_test)           # Evaluate the model on the test da
        LR_visualizer.show()
```



```
Out[ ]: <AxesSubplot:title={'center':'ROC Curves for ExtraTreesClassifier'}, xlabel='False Positive Rate', ylabel='True Positive Rate'>
```

```
In [ ]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred_xtest, cmap='Reds')
```

```
Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1bd5a3a8b50>
```



```
In [ ]: LR_mul = multilabel_confusion_matrix(y_test, y_pred_xtest)
LR_mul
```

```
Out[ ]: array([[6155, 1],
               [ 89, 1096]],

              [[5015, 282],
               [ 279, 1765]],

              [[2888, 341],
               [ 256, 3856]]], dtype=int64)
```

```
In [ ]: LR_FN = LR_mul[2][1][0]
update_score_card(ET_Model, LR_FN, 'Extra Tree')
```

XGBClassifier

```
In [ ]: xgb=XGBClassifier()
XGB_Model=xgb.fit(X_train,y_train)
y_pred_xtest=XGB_Model.predict(X_test)
print(classification_report(y_test,y_pred_xtest))
```


	precision	recall	f1-score	support
0	1.00	0.92	0.96	1185
1	0.88	0.92	0.90	2044
2	0.95	0.95	0.95	4112
accuracy			0.94	7341
macro avg	0.94	0.93	0.94	7341
weighted avg	0.94	0.94	0.94	7341

```
In [ ]: print(accuracy_score(y_test,y_pred_xtest))
        print(confusion_matrix(y_test,y_pred_xtest))
```

```
0.9370657948508377
[[1096  35  54]
 [   0 1887 157]
 [   1  215 3896]]
```

```
In [ ]: LR_visualizer = ROCAUC(XGB_Model)


        LR_visualizer.fit(X_train, y_train)           # Fit the training data to the visu
        LR_visualizer.score(X_test, y_test)           # Evaluate the model on the test da
        LR_visualizer.show()
```



```
Out[ ]: <AxesSubplot:title={'center':'ROC Curves for XGBClassifier'}, xlabel='False P
ositive Rate', ylabel='True Positive Rate'>
```

```
In [ ]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred_xtest, cmap='Reds')
```

```
Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1bd5a4517
60>
```



```
In [ ]: LR_mul = multilabel_confusion_matrix(y_test, y_pred_xtest)
LR_mul
```

```
Out[ ]: array([[[6155, 1],
                [ 89, 1096]],

               [[5047, 250],
                [ 157, 1887]],

               [[3018, 211],
                [ 216, 3896]]], dtype=int64)
```

```
In [ ]: LR_FN = LR_mul[2][1][0]
update_score_card(XGB_Model, LR_FN, 'XGB')
```

LGBMClassifier

```
In [ ]: lgbm=LGBMClassifier()
LGBM_Model=lgbm.fit(X_train,y_train)
y_pred_xtest=LGBM_Model.predict(X_test)
print(classification_report(y_test,y_pred_xtest))
```



	precision	recall	f1-score	support
0	1.00	0.92	0.96	1185
1	0.89	0.92	0.90	2044
2	0.95	0.95	0.95	4112
accuracy			0.94	7341
macro avg	0.94	0.93	0.94	7341
weighted avg	0.94	0.94	0.94	7341

```
In [ ]: print(accuracy_score(y_test,y_pred_xtest))
        print(confusion_matrix(y_test,y_pred_xtest))
```

```
0.9361122462879716
[[1096  33  56]
 [  1 1872 171]
 [  0  208 3904]]
```

```
In [ ]: LR_visualizer = ROCAUC(LGBM_Model)


        LR_visualizer.fit(X_train, y_train)      # Fit the training data to the visu
        LR_visualizer.score(X_test, y_test)      # Evaluate the model on the test da
        LR_visualizer.show()
```



```
Out[ ]: <AxesSubplot:title={'center':'ROC Curves for LGBMClassifier'}, xlabel='False
        Positive Rate', ylabel='True Positive Rate'>
```

```
In [ ]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred_xtest, cmap='Reds')
```

```
Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1bd5a8fab
        80>
```



```
In [ ]: LR_mul = multilabel_confusion_matrix(y_test, y_pred_xtest)
LR_mul
```

```
Out[ ]: array([[[6155,    1],
                [ 89, 1096]],

               [[5056,   241],
                [ 172, 1872]],

               [[3002,   227],
                [ 208, 3904]]], dtype=int64)
```

```
In [ ]: LR_FN = LR_mul[2][1][0]
update_score_card(LGBM_Model, LR_FN, 'LGBM')
```

AdaBoostClassifier

```
In [ ]: ada=AdaBoostClassifier()
ADA_Model=ada.fit(X_train,y_train)
y_pred_xtest=ADA_Model.predict(X_test)
print(classification_report(y_test,y_pred_xtest))
```


	precision	recall	f1-score	support
0	1.00	0.92	0.96	1185
1	0.87	0.91	0.89	2044
2	0.94	0.94	0.94	4112
accuracy			0.93	7341
macro avg	0.94	0.93	0.93	7341
weighted avg	0.93	0.93	0.93	7341

```
In [ ]: print(accuracy_score(y_test,y_pred_xtest))
        print(confusion_matrix(y_test,y_pred_xtest))
```

```
0.9298460700177088
[[1096  33  56]
 [   0 1856 188]
 [   0  238 3874]]
```

```
In [ ]: LR_visualizer = ROCAUC(ADA_Model)


LR_visualizer.fit(X_train, y_train)      # Fit the training data to the visu
LR_visualizer.score(X_test, y_test)      # Evaluate the model on the test da
LR_visualizer.show()
```



```
Out[ ]: <AxesSubplot:title={'center':'ROC Curves for AdaBoostClassifier'}, xlabel='False Positive Rate', ylabel='True Positive Rate'>
```

```
In [ ]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred_xtest, cmap='Reds')
```

```
Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1bd5a3a8a00>
```



```
In [ ]: LR_mul = multilabel_confusion_matrix(y_test, y_pred_xtest)
LR_mul
```

```
Out[ ]: array([[6156,  0],
               [ 89, 1096]],

               [[5026, 271],
                [ 188, 1856]],

               [[2985, 244],
                [ 238, 3874]]], dtype=int64)
```

```
In [ ]: LR_FN = LR_mul[2][1][0]
update_score_card(ADA_Model, LR_FN, 'ADA Boosting')
```

GradientBoostingClassifier

```
In [ ]: gb=GradientBoostingClassifier()
GB_Model=gb.fit(X_train,y_train)
y_pred_xtest=GB_Model.predict(X_test)
print(classification_report(y_test,y_pred_xtest))
```


	precision	recall	f1-score	support
0	1.00	0.92	0.96	1185
1	0.89	0.91	0.90	2044
2	0.94	0.95	0.95	4112
accuracy			0.94	7341
macro avg	0.94	0.93	0.94	7341
weighted avg	0.94	0.94	0.94	7341

```
In [ ]: print(accuracy_score(y_test,y_pred_xtest))
        print(confusion_matrix(y_test,y_pred_xtest))
```

```
0.9367933524043046
[[1096  35  54]
 [  0 1870 174]
 [  0  201 3911]]
```

```
In [ ]: LR_visualizer = ROCAUC(GB_Model)

        LR_visualizer.fit(X_train, y_train)           # Fit the training data to the visu
        LR_visualizer.score(X_test, y_test)           # Evaluate the model on the test da
        LR_visualizer.show()
```



```
Out[ ]: <AxesSubplot:title={'center':'ROC Curves for GradientBoostingClassifier'}, xlabel='False Positive Rate', ylabel='True Positive Rate'>
```

```
In [ ]: LR_mul = multilabel_confusion_matrix(y_test, y_pred_xtest)
        LR_mul
```


```
Out[ ]: array([[[6156,    0],
                [ 89, 1096]],

               [[5061,  236],
                [ 174, 1870]],

               [[3001,  228],
                [ 201, 3911]]], dtype=int64)
```

```
In [ ]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred_xtest, cmap='Reds')
```

```
Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1bd70ea8e50>
```



```
In [ ]: LR_FN = LR_mul[2][1][0]
         update_score_card(GB_Model, LR_FN, 'Gradient Boosting')
```

```
In [ ]: score_card.sort_values(by=['f1-score'], ascending=False)
```

Out[]:

	Model	Precision Score	Recall Score	False Negatives	Kappa Score	f1-score
4	XGB	0.943575	0.931852	216	0.891763	0.937088
7	Gradient Boosting	0.944284	0.930295	201	0.891061	0.936731
5	LGBM	0.943361	0.930054	208	0.889948	0.936146
2	Random Forest	0.940595	0.924845	211	0.882204	0.932166
6	ADA Boosting	0.937779	0.925013	238	0.879255	0.93079
3	Extra Tree	0.926692	0.908714	256	0.852936	0.917194
1	Decision Tree	0.902946	0.904135	290	0.843953	0.903478
0	Logistic Regression	0.671991	0.633901	618	0.464028	0.64833

In []: GB_Model.feature_importances_

Out[]: array([1.73550201e-04, 5.54701251e-04, 1.44876412e-01, 3.87204143e-01,
2.15691768e-05, 2.67363105e-05, 2.15027256e-01, 6.64687808e-06,
5.08928017e-05, 2.11453484e-05, 2.42982402e-05, 4.76905085e-06,
9.81825029e-06, 2.12489040e-05, 1.95503585e-05, 2.51957262e-01])

Tuning Parameters

```
from sklearn.model_selection import GridSearchCV  
params=[{'criterion':['gini','entropy','log_loss'],  
'n_estimators':[100,200,500,1000], 'min_samples_split':[2,4,6,8], 'max_depth':[2,4,6,8]}]  
rf=RandomForestClassifier() grid=GridSearchCV(estimator=rf,cv=5,param_grid=params) grid.fit(X,y)  
grid.best_params_
```

In []:

```
In [ ]: rf=RandomForestClassifier(criterion='gini',max_depth=8,n_estimators=500,min_sa  
RF_Model_ad=rf.fit(X_train,y_train)  
y_pred_xtest=RF_Model_ad.predict(X_test)  
print(classification_report(y_test,y_pred_xtest))
```


	precision	recall	f1-score	support
0	1.00	0.90	0.95	1185
1	0.88	0.90	0.89	2044
2	0.93	0.95	0.94	4112
accuracy			0.93	7341
macro avg	0.94	0.92	0.93	7341
weighted avg	0.93	0.93	0.93	7341

```
In [ ]: print(accuracy_score(y_test,y_pred_xtest))  
print(confusion_matrix(y_test,y_pred_xtest))
```

```
0.9276665304454434
[[1072  33  80]
 [   0 1841 203]
 [   0  215 3897]]
```

```
In [ ]: LR_visualizer = ROCAUC(RF_Model_ad)


LR_visualizer.fit(X_train, y_train)      # Fit the training data to the visu
LR_visualizer.score(X_test, y_test)     # Evaluate the model on the test da
LR_visualizer.show()
```



```
Out[ ]: <AxesSubplot:title={'center': 'ROC Curves for RandomForestClassifier'}, xlabel='False Positive Rate', ylabel='True Positive Rate'>
```

```
In [ ]: ConfusionMatrixDisplay.from_predictions(y_test, y_pred_xtest, cmap='Reds')
```

```
Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1bd6a9db5e0>
```

```
In [ ]: LR_mul = multilabel_confusion_matrix(y_test, y_pred_xtest)
LR_mul
```

```
Out[ ]: array([[6156,  0],
               [ 113, 1072]],

              [[5049,  248],
               [ 203, 1841]],

              [[2946,  283],
               [ 215, 3897]]], dtype=int64)
```

```
In [ ]: LR_FN = LR_mul[2][1][0]
update_score_card(RF_Model_ad, LR_FN, 'Random Forest Tuned')
```

```
In [ ]:
```

```
params = {"n_estimators": [90,100,110,120,130,140,150,200,250], 'learning_rate':
[1.0,0.1,0.01,0.001,0.0001]} clf = AdaBoostClassifier() Grid =
GridSearchCV(clf,param_grid=params,cv=5) Grid.fit(X_train, y_train) Grid.best_params_
```

```
In [ ]: ada=AdaBoostClassifier(learning_rate=1.0,n_estimators=110)
ADA_Model_ad=ada.fit(X_train,y_train)
y_pred_xtest=ADA_Model_ad.predict(X_test)
print(classification_report(y_test,y_pred_xtest))
```


	precision	recall	f1-score	support
0	1.00	0.92	0.96	1185
1	0.90	0.86	0.88	2044
2	0.92	0.96	0.94	4112
accuracy			0.93	7341
macro avg	0.94	0.92	0.93	7341
weighted avg	0.93	0.93	0.93	7341

```
In [ ]: print(accuracy_score(y_test,y_pred_xtest))
        print(confusion_matrix(y_test,y_pred_xtest))
```

```
0.9291649639013758
[[1096  37  52]
 [   1 1766 277]
 [   2  151 3959]]
```

```
In [ ]: LR_visualizer = ROCAUC(ADA_Model_ad)


        LR_visualizer.fit(X_train, y_train)           # Fit the training data to the visu
        LR_visualizer.score(X_test, y_test)           # Evaluate the model on the test da
        LR_visualizer.show()
```



```
Out[ ]: <AxesSubplot:title={'center':'ROC Curves for AdaBoostClassifier'}, xlabel='False Positive Rate', ylabel='True Positive Rate'>
```

```
In [ ]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred_xtest, cmap='Reds')
```

```
Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1bd70ea8100>
```



```
In [ ]: LR_mul = multilabel_confusion_matrix(y_test, y_pred_xtest)
LR_mul
```

```
Out[ ]: array([[[6153,    3],
                [ 89, 1096]],

               [[5109,  188],
                [ 278, 1766]],

               [[2900,  329],
                [ 153, 3959]]], dtype=int64)
```

```
In [ ]: LR_FN = LR_mul[2][1][0]
update_score_card(ADA_Model_ad, LR_FN, 'ADA Boosting Tuned')
```

```
In [ ]: score_card
```


Out[]:

	Model	Precision Score	Recall Score	False Negatives	Kappa Score	f1-score
0	Logistic Regression	0.671991	0.633901	618	0.464028	0.64833
1	Decision Tree	0.902946	0.904135	290	0.843953	0.903478
2	Random Forest	0.940595	0.924845	211	0.882204	0.932166
3	Extra Tree	0.926692	0.908714	256	0.852936	0.917194
4	XGB	0.943575	0.931852	216	0.891763	0.937088
5	LGBM	0.943361	0.930054	208	0.889948	0.936146
6	ADA Boosting	0.937779	0.925013	238	0.879255	0.93079
7	Gradient Boosting	0.944284	0.930295	201	0.891061	0.936731
8	Random Forest Tuned	0.93786	0.91768	215	0.874909	0.926918
9	ADA Boosting Tuned	0.941444	0.917226	153	0.876708	0.928594

Summary: From the above models, we could see that the Gradient Boosting over other models is holding a very good Precision score, Recall Score, Kappa score, f1-score and very less number of False Negative values for the churn risk class of 2. Since, our business involves False Negative values to be costly for our class 2 churn segment, We feel it is important to lower the False Negative of class 2 segment. Since, Gradient Boosting is having ideal scores, it is practically impossible to have those scores, Hence we choose Random forest as our final model.


```
In [ ]: pd.DataFrame(RF_Model.feature_importances_,index = [X.columns]).plot(kind='bar'
```

Out[]: <AxesSubplot:>




```
In [ ]: pd.crosstab(df.membership_category,df.churn_risk_score).plot(kind='bar')
```

```
Out[ ]: <AxesSubplot:xlabel='membership_category'>
```



```
In [ ]: pd.crosstab(df.feedback,df.churn_risk_score).plot(kind='bar')
```

```
Out[ ]: <AxesSubplot:xlabel='feedback'>
```



Business Recommendations

Based on EDA observations and model predictions, the following suggestions have been made:

- Most of the feedbacks are Not Specified, poor customer service, poor website and poor product quality which is affecting our model prediction. The Web development team can make the feedback windows as a compulsory one.
- Quality customer care should be provided. They should address all the concerns.
- Customer with Basic membership and No membership are tend to get churn risk rate 3
- We should try to upsell the memberships. Converting memberships to next level will provide the customers with extra benefits.
- We have to provide offers and bonus for our existing customers. This will improve our revenue in a long run.

In []:

In []:

In []: